


COMP 285

Analysis of Algorithms

Welcome to COMP 285

Lecture 11: BSTs + Self-Balancing Trees

Lecturer: Chris Lucas (cflucas@ncat.edu)



HW2

Grades were released!

HW3

Due Today @ 11:59pm

HW4

Released by EoD! Due 10/13!

Midterm Approaching!

10/06 @ 2pm! Lectures 0-11

Midterm Approaching!

10/06 @ 2pm! Lectures 0-11

Written exam, similar to quizzes, homework assignments, coding

Midterm Approaching!

Review topics?

Written exam, similar to quizzes, homework assignments, coding

Extra Credit Opportunities!

- Technical Interview Prep with Meta ([Oct.](#)) +0.5%
- Technical Interview Prep with Chris ([Wed](#)) +0.5%, up to 1%
 - Midpoint survey ([link](#)) +1%

Quiz!

www.comp285-fall22.ml



**Recall where we
ended last lecture...**

Binary tree terminology

Each node has two **children**.

The **left child** of **3** is **2**

The **right child** of **3** is **4**

The **parent** of **3** is **5**

2 is a **descendant** of **5**


Each node has a pointer to its left child, right child, and parent.

Both **children** of **1** **NIL**.
(I won't usually draw them).

The **height** of this tree is 3. (Max number of edges from the root to a leaf).


This node is the **root**

This is a **node**.
It has a **key** (7).




findMaxOfTree

Instructions: Write an algorithm that takes in a tree of ints, and returns the max of all the values within the tree.



Write an algorithm that takes in a tree of ints, and returns the max of all the values within the tree.




```
int findMaxOfTree(TreeNode<int>* root) {  
    if (root->isLeaf()) {  
        return root->getValue();  
    }  
    int leftMax = INT_MIN;  
    if (root->getLeft()) {  
        leftMax = fMOT(root->getLeft());  
    }  
    int rightMax = INT_MIN;  
    if (root->getRight()) {  
        rightMax = fMOT(root->getRight());  
    }  
    return std::max({  
        root->getValue(),  
        leftMax,  
        rightMax  
    });  
}
```

Big Questions!

- What are/Why binary search trees (BST)?
- Why does balance matter?
- How do we maintain balance?



Our Guiding Questions...



Motivation for Binary Search Trees

| | Sorted Arrays | Linked Lists | (balanced) Binary Search Trees |
|--------|---------------|--------------|--------------------------------------|
| Search | | | |
| Delete | | | |
| Insert | | | |

Motivation for Binary Search Trees

| | Sorted Arrays | Linked Lists | (balanced) Binary Search Trees |
|--------|---------------|--------------|--------------------------------------|
| Search | $O(\log(n))$ | | |
| Delete | $O(n)$ | | |
| Insert | $O(n)$ | | |

Motivation for Binary Search Trees


| | Sorted Arrays | Linked Lists | (balanced) Binary Search Trees |
|--------|---------------|--------------|--------------------------------------|
| Search | $O(\log(n))$ | $O(n)$ | |
| Delete | $O(n)$ | $O(n)$ | |
| Insert | $O(n)$ | $O(1)$ | |

Motivation for Binary Search Trees

| | Sorted Arrays | Linked Lists | (balanced) Binary Search Trees |
|--------|---------------|--------------|--------------------------------------|
| Search | $O(\log(n))$ | $O(n)$ | $O(\log(n))$ |
| Delete | $O(n)$ | $O(n)$ | $O(\log(n))$ |
| Insert | $O(n)$ | $O(1)$ | $O(\log(n))$ |


Binary Search Trees

- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:




Binary Search Trees

- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:




Binary Search Trees

- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:




Binary Search Trees

- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:



Binary Search Trees

- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.
- Example of building a binary search tree:




Q: Is this the only binary search tree I could possibly build with these values?

A: **No.** I made choices about which nodes to choose when. Any choices would have been fine.



Aside: this should look familiar

kinda like QuickSort



Are these Binary Search Trees? Yes or No?


- A BST is a binary tree so that:
 - Every LEFT descendant of a node has key less than that node.
 - Every RIGHT descendant of a node has key larger than that node.




Are these Binary Search Trees? Yes or No?

- A BST is a binary tree so that:
 - Every LEFT descendant of a node is less than the node
 - Every RIGHT descendant of a node is greater than the node


Can we do this programmatically?





**Given a root node,
determine if it's a BST
and return True/False**



Given a root node,
determine if it's a BST
and return True/False



Given a root node,
determine if it's a BST and
return True/False




**Let's code
it!!!**




Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`




Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if $x \neq \text{NIL}$:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`




Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`




Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`




Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`




Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`



Aside: In-Order Traversal of BSTs


- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`



2

Aside: In-Order Traversal of BSTs


- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if $x \neq \text{NIL}$:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`



2 3

Aside: In-Order Traversal of BSTs


- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`



2 3 4

Aside: In-Order Traversal of BSTs


- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`



2 3 4

Aside: In-Order Traversal of BSTs


- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`



2 3 4 5

Aside: In-Order Traversal of BSTs


- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if `x != NIL`:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`



2 3 4 5 7

Aside: In-Order Traversal of BSTs

- Output all the elements in sorted order!
- `inOrderTraversal(x)`:
 - if $x \neq \text{NIL}$:
 - `inOrderTraversal(x.left)`
 - `print(x.value)`
 - `inOrderTraversal(x.right)`
- Runs in time $O(n)$.



2 3 4 5 7 Sorted!

Big Questions!

- What are/Why binary search trees?
- Why does balance matter?
- How do we maintain balance?




Back to the goal

Fast SEARCH/INSERT/DELETE

Can we do these?

| | Sorted Arrays | Linked Lists | Binary Search Trees* |
|--------|----------------|--------------|----------------------|
| Search | $O(\log(n))$ 😊 | $O(n)$ 😞 | $O(\log(n))$ 😊 |
| Delete | $O(n)$ 😞 | $O(n)$ 😞 | $O(\log(n))$ 😊 |
| Insert | $O(n)$ 😞 | $O(1)$ 😊 | $O(\log(n))$ 😊 |

SEARCH in a Binary Search Tree



EXAMPLE: Search for 4.


EXAMPLE: Search for 4.5

- It turns out it will be convenient to **return 4** in this case
- (that is, **return** the last node before we went off the tree)

How long does this take?

$O(\text{length of longest path}) = O(\text{height})$

SEARCH in a Binary Search Tree



EXAMPLE: Search for 4.

EXAMPLE: Search for 4.5


- **if** value > x.value:
 - Recurse on x.rightChild
- **if** value < x.value:
 - Recurse on x.leftChild
- **if** x.value == value:
 - **Return x**

How long does this take?

$O(\text{length of longest path}) = O(\text{height})$


INSERT in a Binary Search Tree

EXAMPLE: Insert 4.5



- **INSERT**(value):
 - $x = \text{SEARCH}(\text{value})$
 - **Insert** a new node with desired value at x ...

INSERT in a Binary Search Tree




EXAMPLE: Insert 4.5

- **INSERT(value):**
 - $x = \text{SEARCH}(\text{value})$
 - **if** $\text{value} > x.\text{value}$:
 - Make a new node with the correct value, and put it as the right child of x .
 - **if** $\text{value} < x.\text{value}$:
 - Make a new node with the correct value and put it as the left child of x .
 - **if** $x.\text{value} == \text{value}$:
 - **return**

DELETE in a Binary Search Tree


EXAMPLE: Delete 2




- **DELETE**(value):
 - $x = \text{SEARCH}(\text{value})$
 - **if** $x.\text{value} == \text{value}$:
 -delete x

DELETE in a Binary Search Tree

say we want to delete 3




Case 1: if 3 is a leaf,
just delete it.



Case 2: if 3 has just one child,
move that up.

DELETE in a Binary Search Tree (cont.)


Case 3: if 3 has two children, replace 3 with its **immediate successor**.
(aka, next biggest thing after 3)



- Does this maintain the BST property?
 - Yes.
- How do we find the immediate successor?
 - SEARCH for 3 in the subtree under 3.right
- How do we remove it when we find it?
 - If [3.1] has 0 or 1 children, do one of the previous cases.
- What if [3.1] has two children?
 - It doesn't.

How long do these operations take?

- **SEARCH** is the big one.
 - Everything else just calls **SEARCH** and then does some small $O(1)$ -time operation.



Time = $O(\text{height of tree})$

Trees have depth $O(\log(n))$. **Done!**

Wait a second...




How long does search take?

1 minute think; 1 minute pair+share



Search might take time $O(n)$.



- This is a valid binary search tree.
- The version with n nodes has depth n , **not** $O(\log(n))$.

What to do?


- Goal: Fast SEARCH/INSERT/DELETE
- All these things take time $O(\text{height})$
- And the height might be big!!! 😞
- Idea!
 - Keep track of how deep the tree is getting.
 - If it gets too tall, re-do everything from scratch.
 - At least $\Omega(n)$ every so often....
- Turns out that's not a great idea, but we're onto something...

Big Questions!

- What are/Why binary search trees?
- Why does balance matter?
- How do we maintain balance?




Self-Balancing Binary Search Trees



Idea! Rotations

- Maintain Binary Search Tree (BST) property, while moving stuff around.

Note: A, B, C, X, Y are
variable names, not the
contents of the nodes.





Idea! Rotations

- Maintain Binary Search Tree (BST) property, while moving stuff around.



Note: A, B, C, X, Y are variable names, not the contents of the nodes.

YOINK!



CLAIM:
this still has BST property.

This seems helpful



Next idea! Have some proxy for balance

- Maintaining perfect balance is too hard.
- Instead, come up with some proxy for balance:
 - If the tree satisfies [SOME PROPERTY], then it's pretty balanced.
 - We can maintain [SOME PROPERTY] using rotations.




There are actually several ways to do this, but today we'll see...

Red-Black Trees

- A Binary Search Tree that balances itself!
- No more time-consuming by-hand balancing!

Red-Black tree!


*Maintain balance by stipulating that black nodes are balanced, and that there *aren't too many* red nodes.*




Red-Black Trees obey the following rules (which are a proxy for balance)

- Every node is colored **red** or **black**.
- The root node is a **black** node.
- NIL children count as **black** nodes.
- Children of a **red** node are **black** nodes.
- For all nodes x :
 - all paths from x to NIL's have the same number of **black** nodes on them.

I'm not going to draw the NIL children in the future, but they are treated as black nodes.






Examples(?)



Which of these
are red-black trees?
(NIL nodes not drawn)


Yes!

- Every node is colored **red** or **black**.
- The root node is a **black node**.
- NIL children count as **black nodes**.
- Children of a **red node** are **black nodes**.
- For all nodes x :
 - all paths from x to NIL's have the same number of **black nodes** on them.



Why these rules???????

- This is pretty balanced.
 - The **black nodes** are balanced
 - The **red nodes** are “spread out” so they don’t mess things up too much.
- We can maintain this property as we insert/delete nodes, by using rotations.



This is the really clever idea!

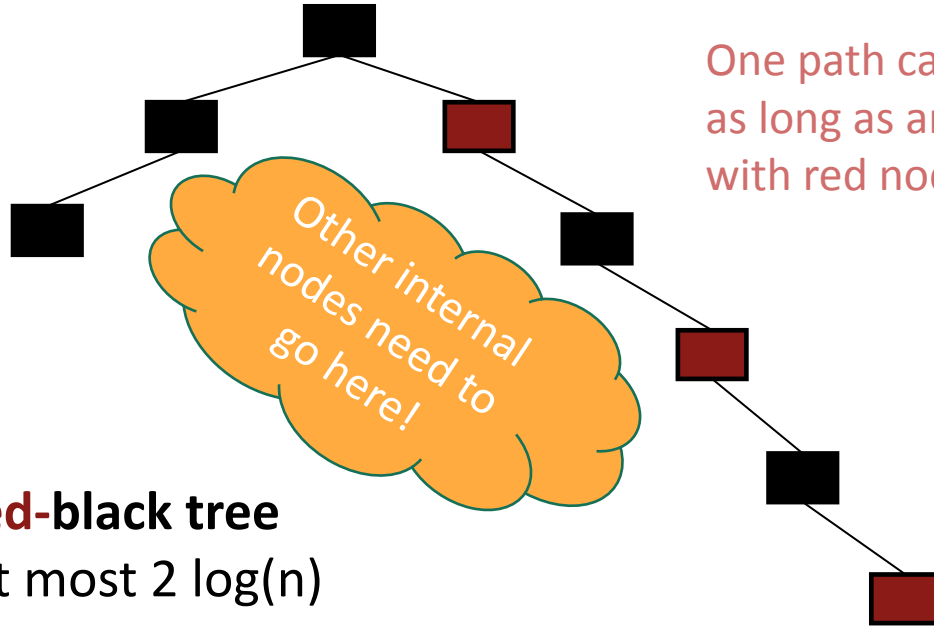
This **Red-Black** structure is a **proxy for balance**.

It’s just a smidge weaker than perfect balance, but we can actually maintain it!

This is “pretty balanced”

- To see why, intuitively, let's try to build a Red-Black Tree that's unbalanced.

Note, this is just a conjecture to build intuition! Rigorous proof on the next, skipped slide.



One path can be at most twice as long as another if we pad it with red nodes.

Conjecture:
the height of a **red-black tree**
with n nodes is at most $2 \log(n)$



RBTree with N non-nil nodes

- Define $b(x)$ to be the number of black nodes in any path from x to NIL.
 - (excluding x , including NIL).

- Claim:

- There are at least $2^{b(x)} - 1$ non-NIL nodes in the subtree underneath x . (Including x).

- [You can proof by induction]

Then:

$$n \geq 2^{b(\text{root})} - 1$$


using the Claim

$$\geq 2^{\text{height}/2} - 1$$

$b(\text{root}) \geq \text{height}/2$ because of RBTree rules.

Rearranging:

$$n + 1 \geq 2^{\text{height}/2} \Rightarrow \text{height} \leq 2\log(n + 1)$$



Claim: at least $2^{b(x)} - 1$ nodes in this WHOLE subtree (of any color).


This is great!

- SEARCH in an RBTree is immediately $O(\log(n))$, since the depth of an RBTree is $O(\log(n))$.
- What about INSERT/DELETE?
 - Turns out, you can INSERT and DELETE items from an RBTree in time $O(\log(n))$, while *maintaining* the RBTree property.
 - That's why this is a good property!

INSERT/DELETE

- I expect we are out of time...
 - There are some slides which you can check out to see how to do INSERT/DELETE in RBTrees if you are curious.
 - See CLRS Ch 13. for even more details.
- You are **not responsible** for the details of INSERT/DELETE for RBTrees for this class.
 - You should know what the “proxy for balance” property is and why it ensures approximate balance.
 - You should know **that** this property can be efficiently maintained, but you do not need to know the details of how.



INSERT: Many cases



- Suppose we want to insert 0 **here**.
- There are 3 “important” cases for different colorings of the existing tree, and there are 9 more cases for all of the various symmetries of these 3 cases.

INSERT: Case 1

- Make a new **red node**.
- Insert it as you would normally.




What if it looks like this?

Example: insert 0





INSERT: Many cases



- Suppose we want to insert 0 **here**.
- There are 3 “important” cases for different colorings of the existing tree, and there are 9 more cases for all of the various symmetries of these 3 cases.


INSERT: Case 2

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.



What if it looks like this?



Example: insert 0



No!

INSERT: Case 2


- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.



What if it looks like this?


Example: insert 0

Can't we just insert 0 as a **black node**?




No!

We need a bit more context




What if it looks like this?
Example: insert 0




We need a bit more context

- Add 0 as a red node.





What if it looks like this?
Example: insert 0



We need a bit more context




- Add 0 as a red node.
- **Claim:** RB-Tree properties still hold.



What if it looks like this?
Example: insert 0



Flip
colors!

But what if **that** was red?





What if it looks like this?
Example: insert 0

More context...



What if it looks like this?
Example: insert 0

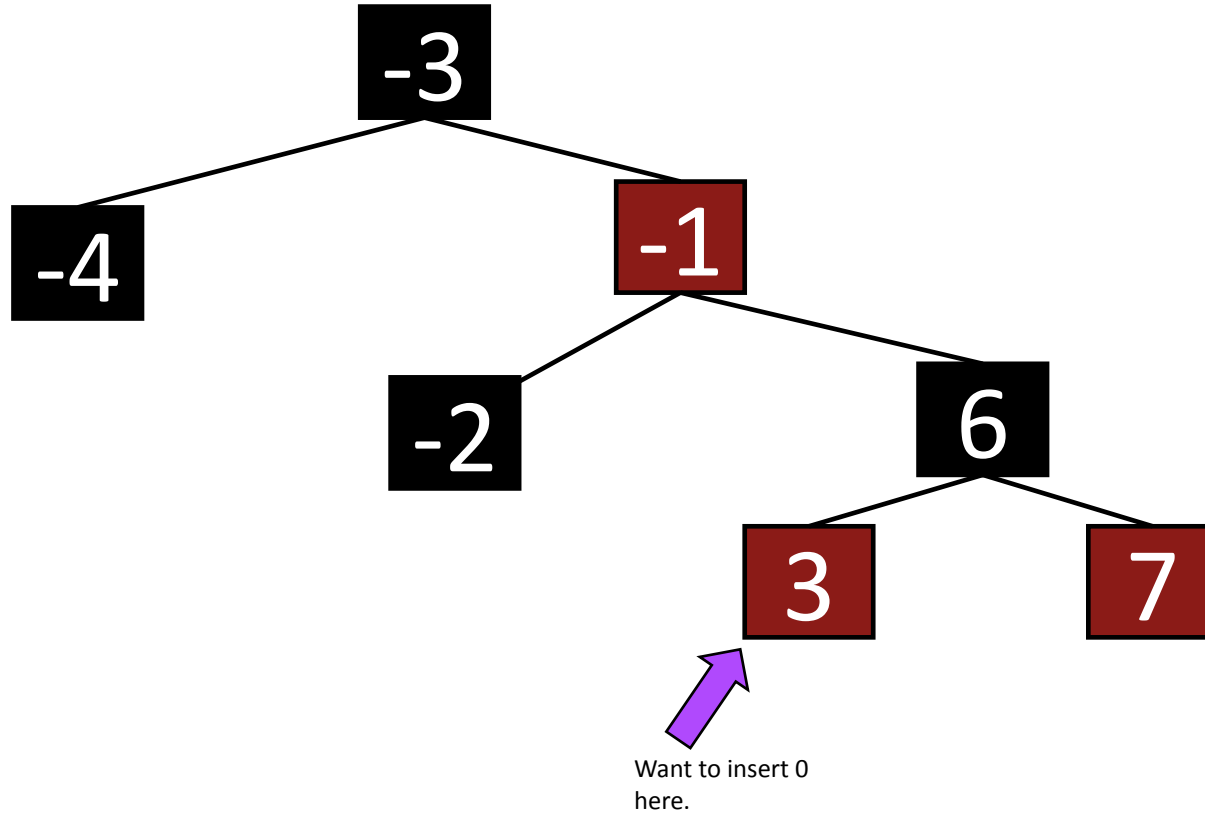
More context...




What if it looks like this?
Example: insert 0

Now we're basically
inserting 6 into some
smaller tree. Recurse!
This one!


Example, part I




Example, part I




Example, part I




Example, part I






Need to know how to insert into trees that look like this...



Want to insert 6 here.





INSERT: Many cases



- Suppose we want to insert 0 **here**.
- There are 3 “important” cases for different colorings of the existing tree, and there are 9 more cases for all of the various symmetries of these 3 cases.

INSERT: Case 3

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.




What if it looks like this?

Example: Insert 0.

- Maybe with a subtree below it.

Recall Rotations



- Maintain Binary Search Tree (BST) property, while moving stuff around.




Inserting into a Red-Black Tree

- Make a new **red node**.
- Insert it as you would normally.
- Fix things up if needed.

YOINK!




What if it looks like this?





Argue that this is a good thing to do!


Example, part 2



Example, part 2




Example, part 2 **YOINK!**




Example, part 2

TA-DA!




Many cases



- Suppose we want to insert 0 **here**.
- There are 3 “important” cases for different colorings of the existing tree, and there are 9 more cases for all of the various symmetries of these 3 cases.

Deleting from a Red-Black tree

Fun exercise!



That's a lot of cases!

- You are **not responsible** for the nitty-gritty details of Red-Black Trees. (For this class)
 - Though implementing them is a great exercise!
- You should know:
 - What are the properties of an RB tree?
 - And (more important) why does that guarantee that they are balanced?

What have we learned?


- Red-Black Trees always have height at most $2\log(n+1)$.
- As with general Binary Search Trees, all operations are $O(\text{height})$
- So all operations with RBTrees are $O(\log(n))$.

Conclusion: The best of both worlds

| | Sorted Arrays | Linked Lists | Binary Search Trees* |
|--------|----------------|--------------|----------------------|
| Search | $O(\log(n))$ 😊 | $O(n)$ 😞 | $O(\log(n))$ 😊 |
| Delete | $O(n)$ 😞 | $O(n)$ 😞 | $O(\log(n))$ 😊 |
| Insert | $O(n)$ 😞 | $O(1)$ 😊 | $O(\log(n))$ 😊 |

Recap

- Balanced binary trees are the best of both worlds!
- But we need to keep them balanced.
- **Red-Black Trees** do that for us.
 - We get $O(\log(n))$ -time INSERT/DELETE/SEARCH
 - Clever idea: have a proxy for balance



How was the pace today?

COMP 285

Analysis of Algorithms

Welcome to COMP 285

Lecture 11: BSTs + Self-Balancing Trees

Lecturer: Chris Lucas (cflucas@ncat.edu)

