


COMP - 285

Analysis of Algorithms

Welcome to COMP 285

Lecture 16: Weighted Graphs, Dijkstra's, A*

Lecturer: Chris Lucas (cflucas@ncat.edu)



Office Hours

Today only -> 3:30pm-4:00pm

HW5 Due!

Tuesday 10/25 @ 11:59PM

HW5 Due!

Q2 (Alphabet Reconstruction) is optional!

HW4

**Recall where we
ended last lecture...**

Analogy

BFS



DFS



Breadth/Depth-First Search Pseudocode

algorithm BFS

Input: undirected graph $G = (V, E)$, s and d

Output: true/false if path from s to d

```
frontier = Queue of integers
visited = {} // empty hash set
frontier.add(s)
visited.insert(s)
while not frontier.empty()
    currNode = frontier.remove()
    if currNode == d
        return true
    for each neighbor of currNode
        if neighbor not in visited
            visited.insert(neighbor)
            frontier.add(neighbor)
return false
```

algorithm DFS

Input: undirected graph $G = (V, E)$, s and d

Output: true/false if path from s to d

```
frontier = Stack of integers
visited = {} // empty hash set
frontier.add(s)
visited.insert(s)
while not frontier.empty()
    currNode = frontier.remove()
    if currNode == d
        return true
    for each neighbor of currNode
        if neighbor not in visited
            visited.insert(neighbor)
            frontier.add(neighbor)
return false
```


Depth-First Search Pseudocode

algorithm DFS

Input: undirected graph G , int s and int t
Output: whether or not there's a path from s to t

```
visited = new boolean array of size  $|V|$   
return dfsHelper( $G$ ,  $s$ ,  $t$ , visited)
```

algorithm DFSHelper

Input: undirected graph $G = (V, E)$, s , t , and visited
Output: whether or not there's a path from s to t

```
visited[s] = true  
if  $s == t$   
    return true  
for each neighbor of  $s$ , starting from smallest  
labeled neighbor  
    if !visited[neighbor] and dfsHelper( $G$ ,  
neighbor,  $t$ , visited)  
        return true  
return false
```

Time Complexity? $O(V + E)$
Space Complexity? $O(V)$

algorithm DFS

Input: undirected graph $G = (V, E)$, s and d
Output: true/false if path from s to d

```
frontier = Stack of integers  
visited = {} // empty hash set  
frontier.add( $s$ )  
visited.insert( $s$ )  
while not frontier.empty()  
    currNode = frontier.remove()  
    if  $currNode == d$   
        return true  
    for each neighbor of currNode  
        if neighbor not in visited  
            visited.insert(neighbor)  
            frontier.add(neighbor)  
return false
```

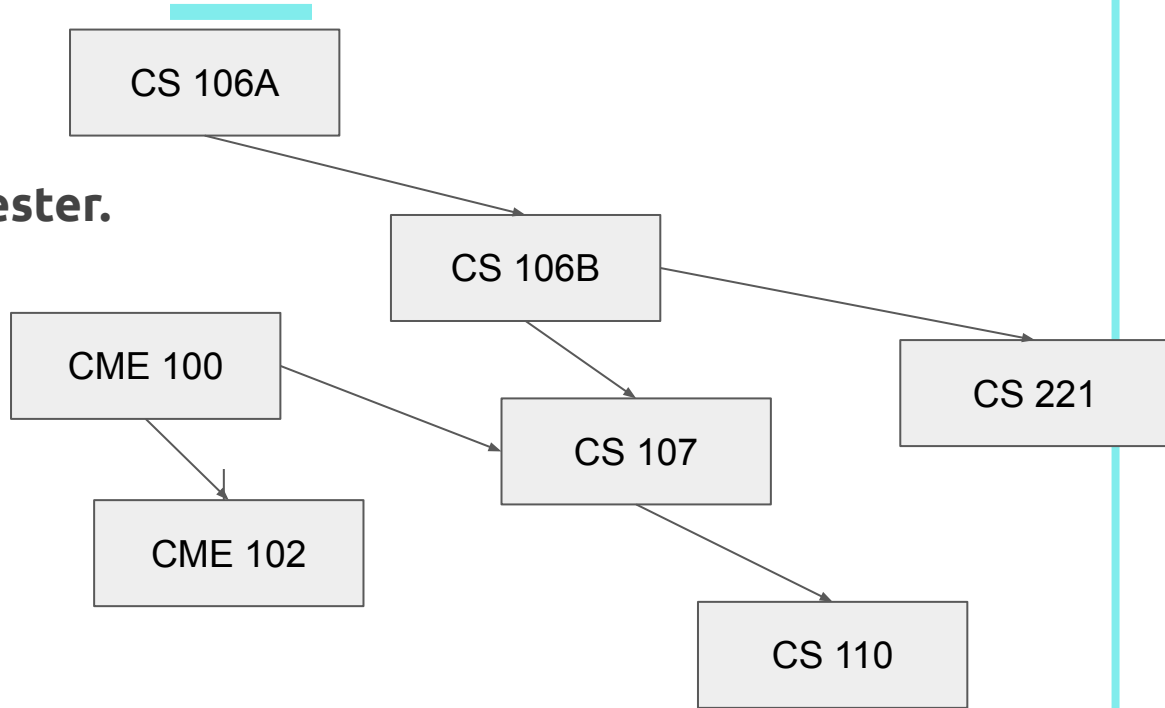
Time Complexity? $O(V+E)$
Space Complexity? $O(V)$

Topological Sort Example #1

What's one valid course ordering?

Assume 1 course per semester.

CME 100
CME 102
CS 106A
CS 106B
CS 107
CS 221
CS 110




Big Questions!

- How do we sort DAGs?
- What's a weighted graph?
- How can we traverse weighted graphs?



Big Questions!




- How do we sort DAGs?
- What's a weighted graph?
- How can we traverse weighted graphs?



Topological Sort Algorithm


- Note that some valid orderings laid out each “child” until we reached a sink.
- We can actually perform a DFS, starting at source nodes and reverse the DFS finishing times to get a valid topological sort.

TopoSort Example



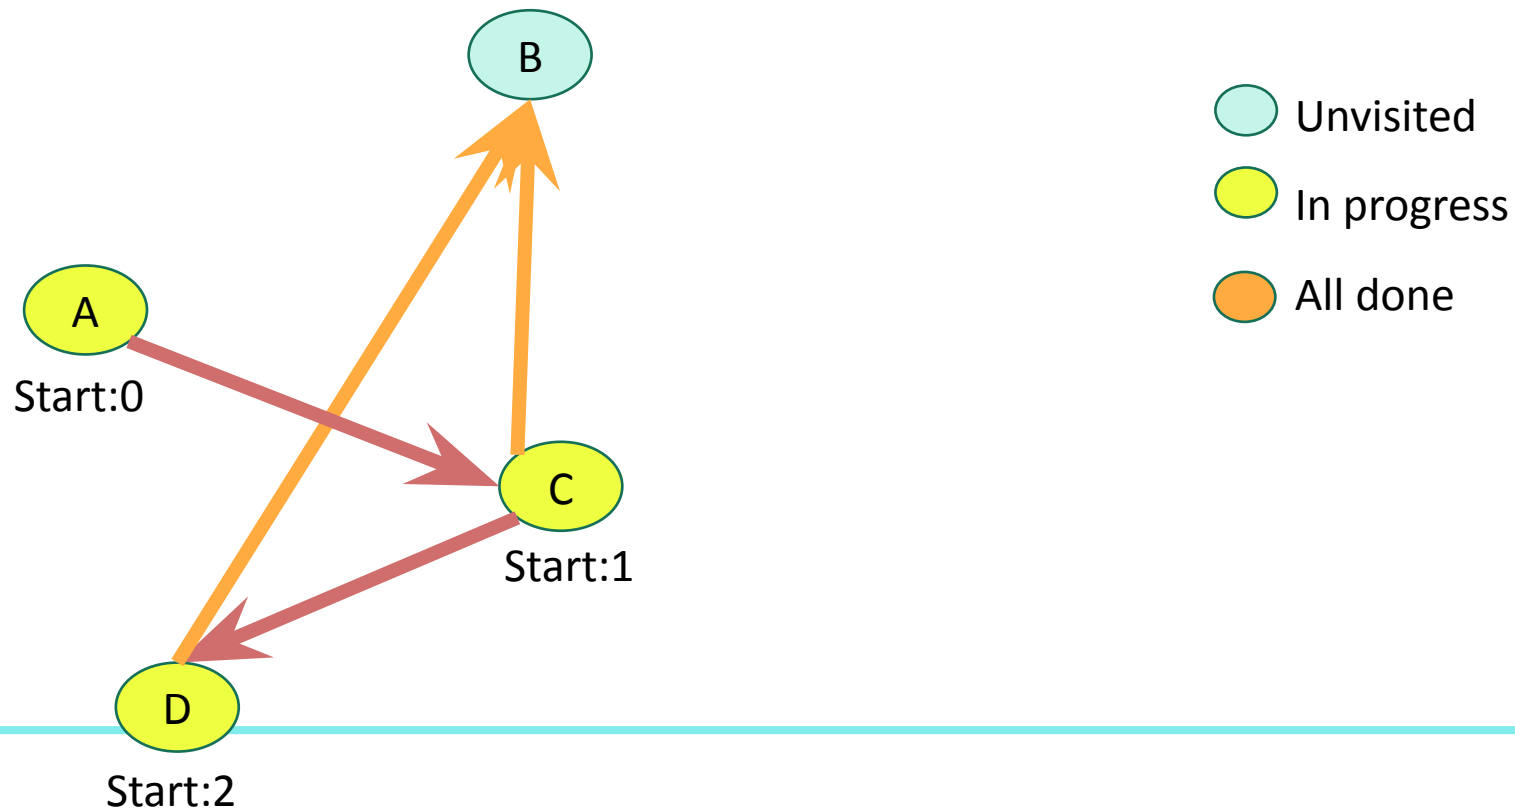
- Unvisited
- In progress
- All done

TopoSort Example

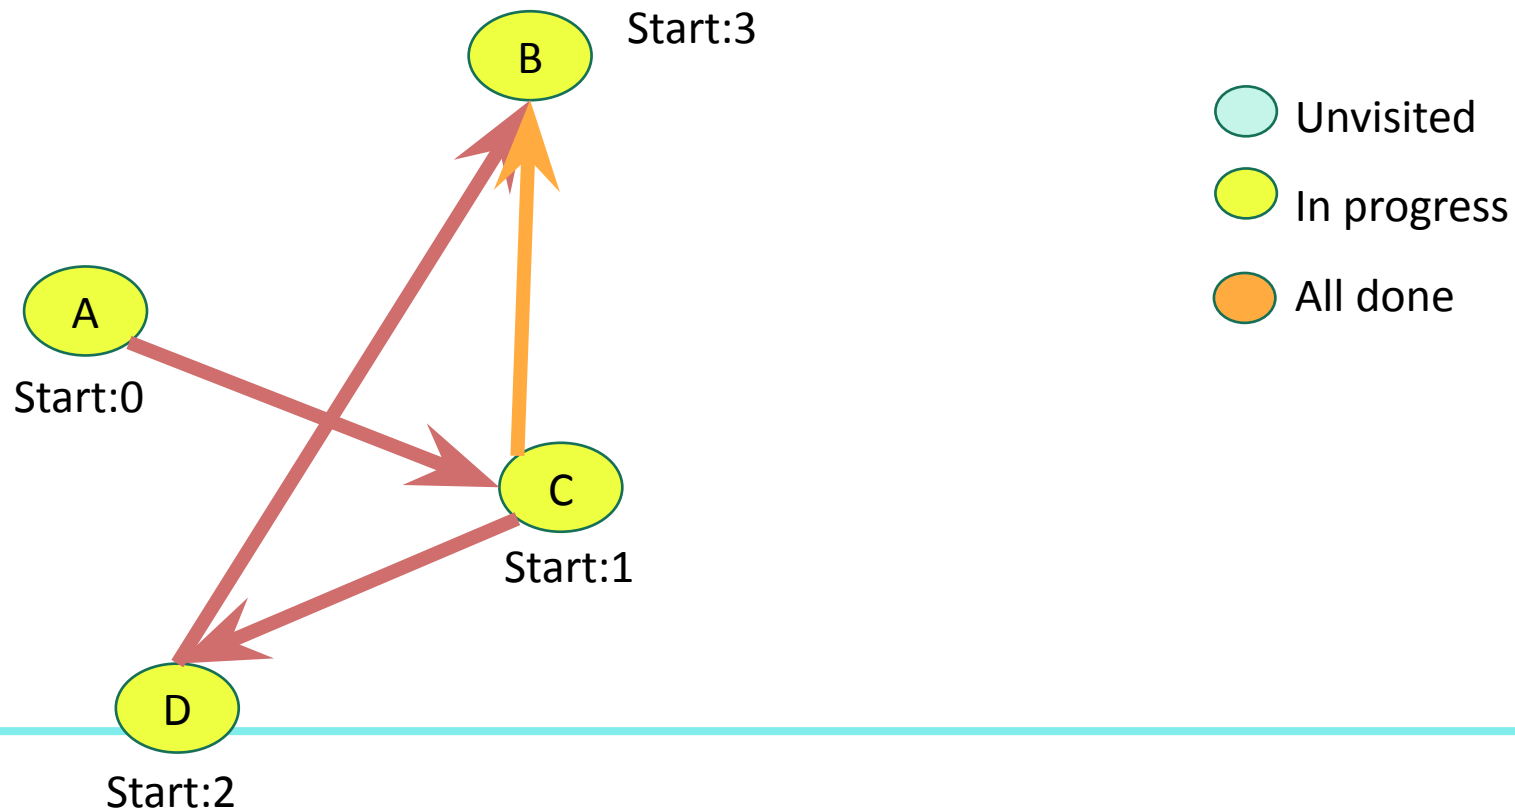


- Unvisited
- In progress
- All done

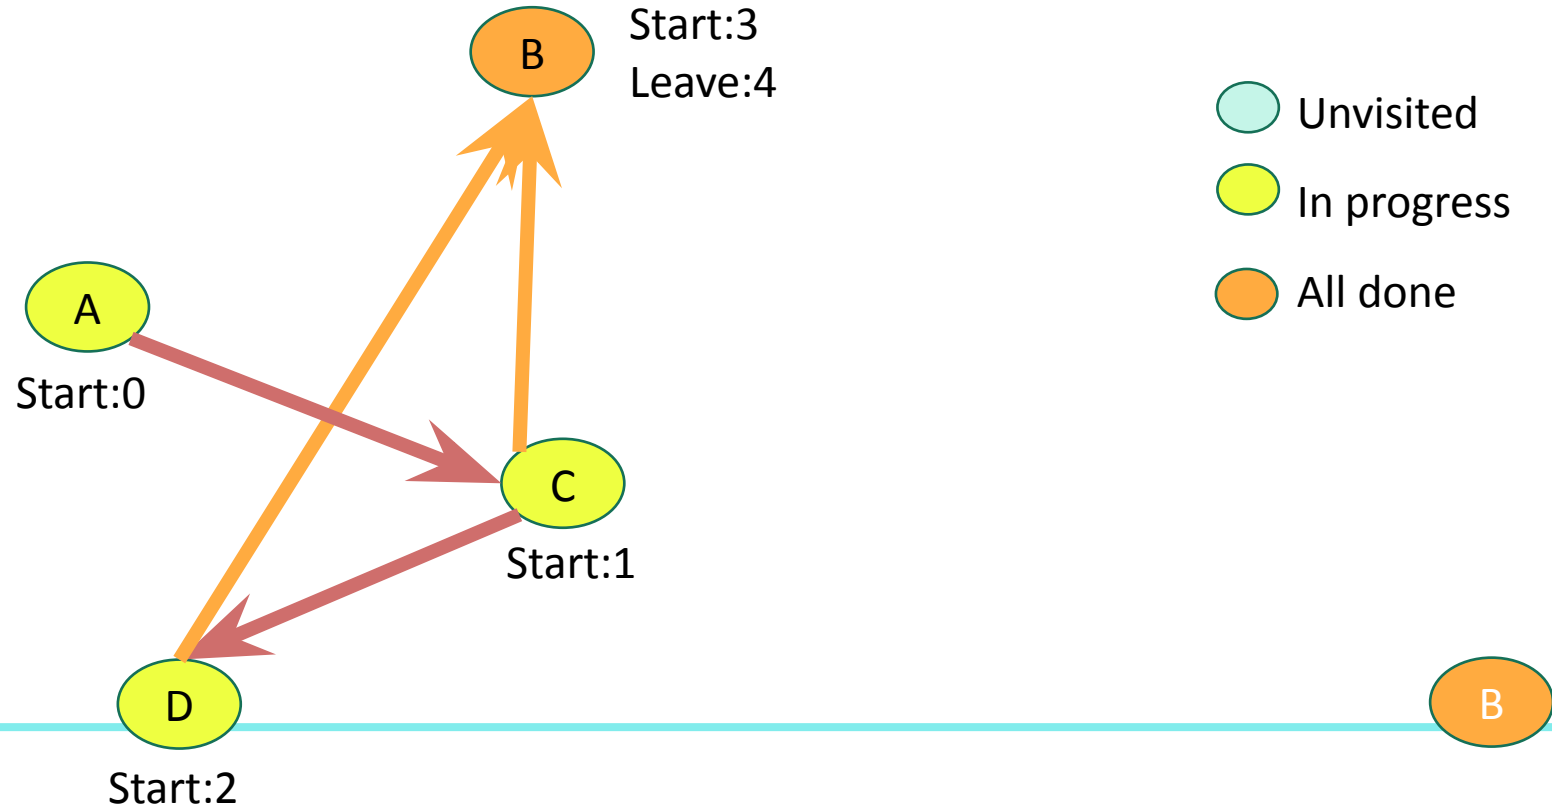
TopoSort Example




TopoSort Example




TopoSort Example




TopoSort Example



- Unvisited
- In progress
- All done




TopoSort Example



- Unvisited
- In progress
- All done




TopoSort Example



- Unvisited
- In progress
- All done

Do them in this order:



TopoSort Pseudocode

algorithm TopoSort

Input: undirected graph G , int s and int t

Output: whether or not there's a path from s to t

visited = new boolean array of size $|V|$

current_label = $|V|$ // global variable

For v in G :

 if !visited[v]:

 return topoSortHelper(G , v , visited)

algorithm topoSortHelper

Input: undirected graph G , v , and visited

visited[v] = true


for each neighbor of v :

 if !visited[neighbor]:

 topoSortHelper(G , neighbor, visited)

$v.\text{end} = \text{current_label}$

 current_label -= 1



TopoSort Pseudocode

algorithm TopoSort

Input: undirected graph G , int s and int t

Output: whether or not there's a path from s to t

visited = new boolean array of size $|V|$

current_label = $|V|$ // global variable

For v in G :

 if !visited[v]:

 return topoSortHelper(G , v , visited)

algorithm topoSortHelper

Input: undirected graph G , v , and visited

visited[v] = true


for each neighbor of v :

 if !visited[neighbor]:


 topoSortHelper(G , neighbor, visited)

v .end = current_label

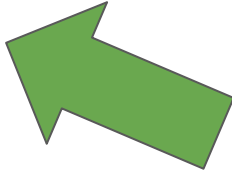
current_label -= 1



see repl.it



Big Questions!

- How do we sort DAGs?
- What's a weighted graph? 
- How can we traverse weighted graphs?



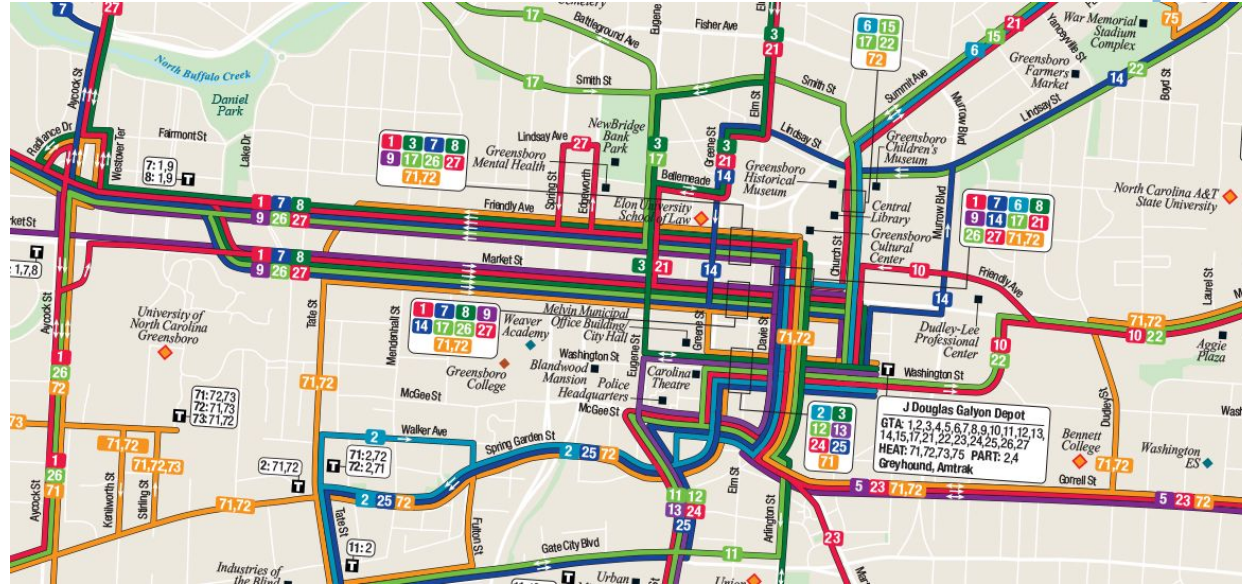
Single-source shortest-path problem

- I want to know the shortest path from one vertex (Aggie Village) to all other vertices.

Destination	Cost	To get there
Clock Tower	1	Clock Tower
Health Center	2	Clock Tower-Health Center
Morrison Hall	10	Morrison Hall
COMP 285	17	COMP 285
Aggie Suites	6	Clock Tower-Health Center-Aggie Suites
Stadium	15	Stadium
Aggie Terrace	21	Clock Tower-Aggie Terrace

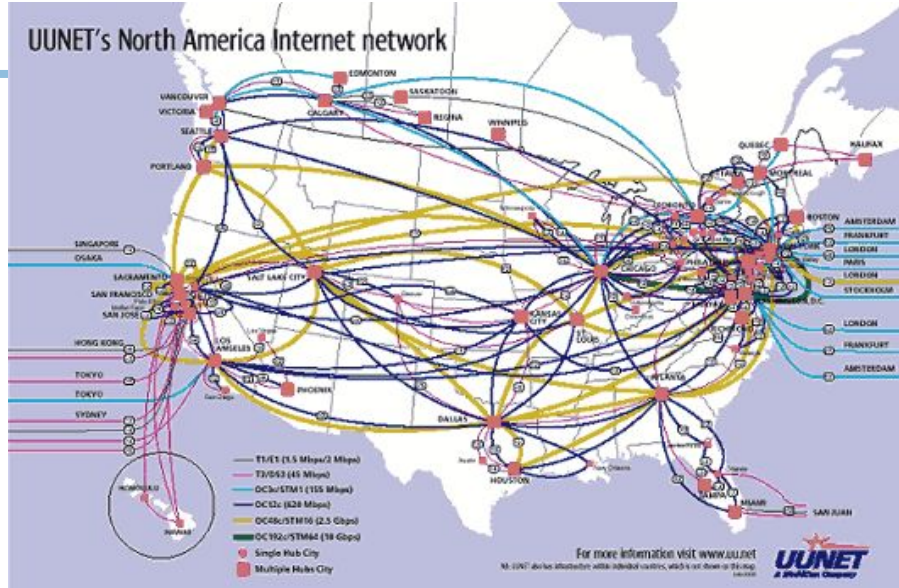
Example

- “what is the shortest path from NC A&T to [anywhere else]” using bus, bus, bike, walking, uber/lyft, driving
- Edge weights have something to do with time, money, hassle.



Example pt. 2

- Network routing
- I send information over the internet, from my computer to all over the world.
- Each path has a cost which depends on link length, traffic, other costs, etc..
- How should we send packets?



```
..rithms-course      traceroute -a www.ethz...  ..lopment/site2      ../hw-solutions


> ~ traceroute -a www.ethz.ch
traceroute to www.ethz.ch (129.132.19.216), 64 hops max, 52 byte packets
 1 [AS0] 192.168.87.1 (192.168.87.1) 3.824 ms 2.489 ms 2.246 ms
 2 [AS0] 192.168.1.254 (192.168.1.254) 4.706 ms 2.512 ms 5.135 ms
 3 [AS0] 104.188.160.1 (104.188.160.1) 5.180 ms 4.191 ms 4.441 ms
 4 [AS0] 99.94.204.178 (99.94.204.178) 3.357 ms 3.007 ms 6.353 ms
 5 [AS7018] 12.242.111.38 (12.242.111.38) 16.523 ms 18.199 ms 14.701 ms
 6 [AS7018] attga402igs.ip.att.net (12.122.117.121) 14.785 ms 13.600 ms 14.220 ms
 7 [AS3356] 4.68.62.225 (4.68.62.225) 10.973 ms 13.979 ms 12.101 ms
 8 [AS3356] ae1.10.bar1.geneva1.level3.net (4.69.203.66) 107.572 ms 107.457 ms 113.640 ms
 9 [AS9057] dante.bar1.geneva1.level3.net (213.242.73.74) 108.367 ms 109.289 ms 109.347 ms
10 [AS559] swice4-b4.switch.ch (130.59.36.70) 109.503 ms 109.147 ms 108.886 ms
11 [AS559] swibf1-b2.switch.ch (130.59.36.113) 125.068 ms 125.046 ms 123.263 ms
12 [AS559] swiez3-b5.switch.ch (130.59.37.6) 121.484 ms 120.742 ms 117.935 ms
13 [AS559] rou-gw-lee-tengig-to-switch.ethz.ch (192.33.92.1) 124.010 ms 124.115 ms 123.322
14 [AS559] rou-fw-rz-rz-gw.ethz.ch (192.33.92.169) 119.290 ms 121.185 ms *
```

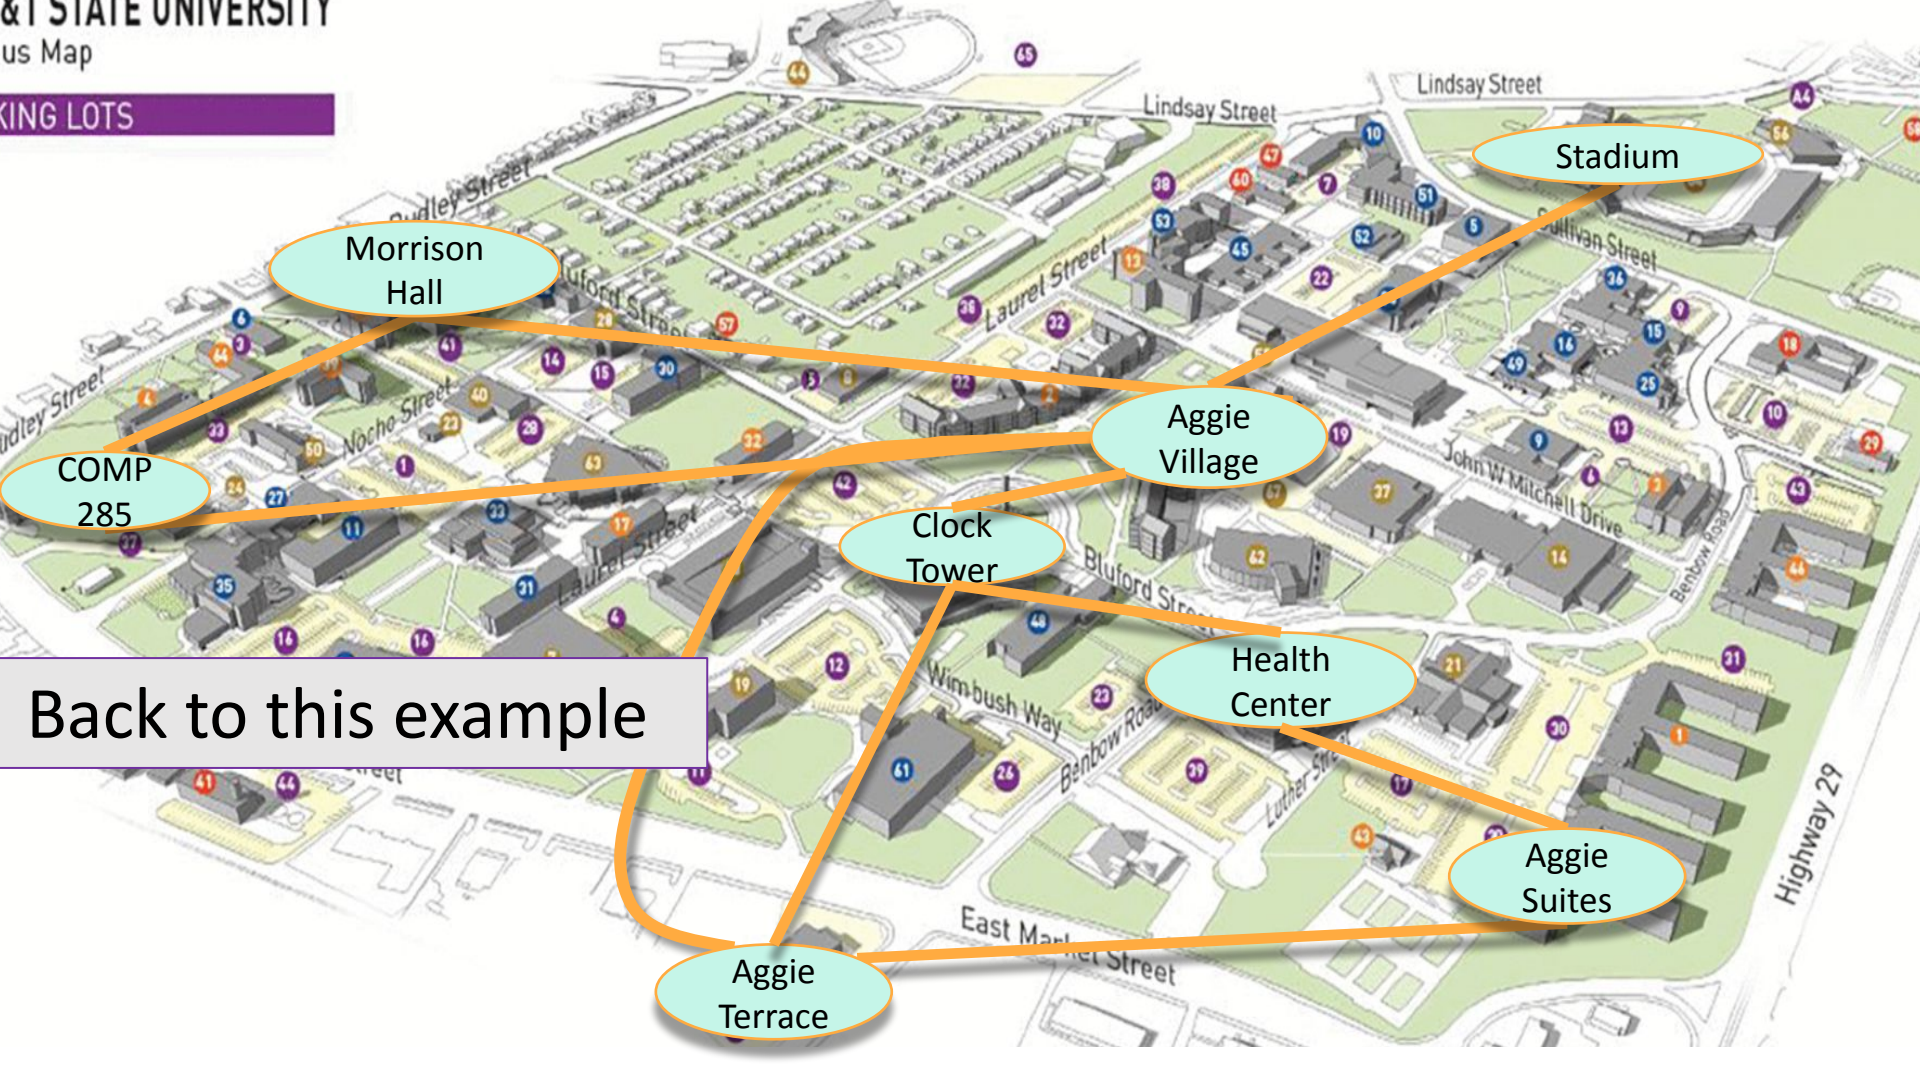
These are difficult, real-world problems

- Costs may change
 - If it's raining, the cost of biking is higher
 - If a link is congested, the cost of routing a packet along it is higher
- The network might not be known
 - My computer doesn't store a map of the internet
- We want to do these tasks really quickly
 - Case and point: **the internet.**

Big Questions!

- How do we sort DAGs?
- What's a weighted graph?
- How can we traverse weighted graphs?





Morrison
Hall

COMP
285

Stadium

Aggie
Village

Clock
Tower

Health
Center


Aggie
Suites

Aggie
Terrace

Back to this example

Dijkstra's algorithm


- Finds shortest paths from Aggie Village to everywhere else.



Dijkstra

intuition

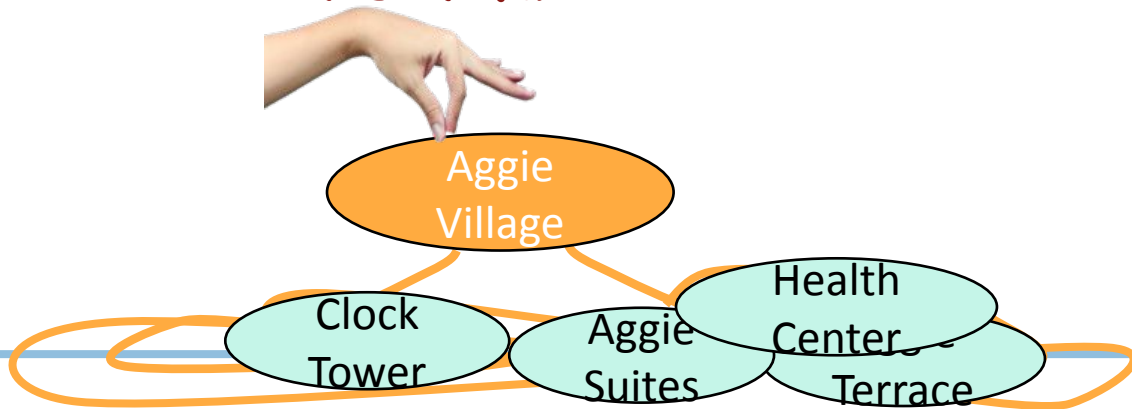
YOINK!



Dijkstra intuition


A vertex is done when it's not on the ground anymore.

YOINK!




Dijkstra intuition

YOINK!




Dijkstra intuition

YOINK!



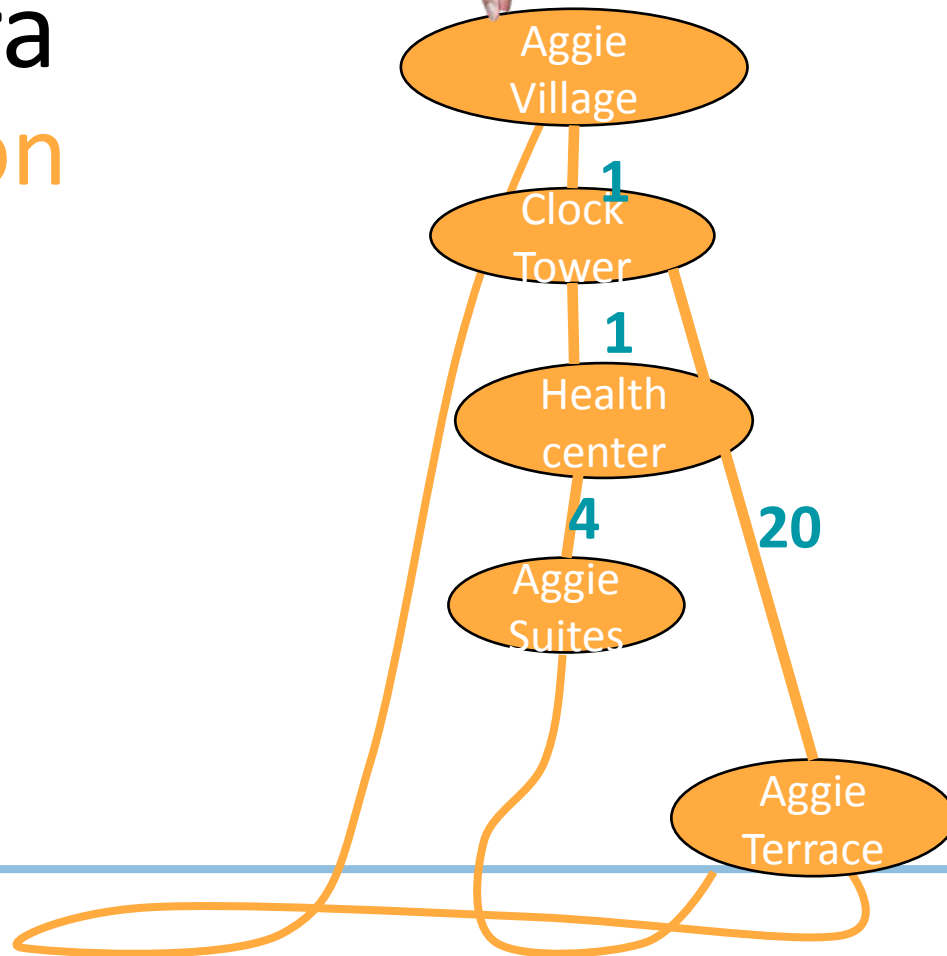
Dijkstra intuition

YOINK!



Dijkstra intuition

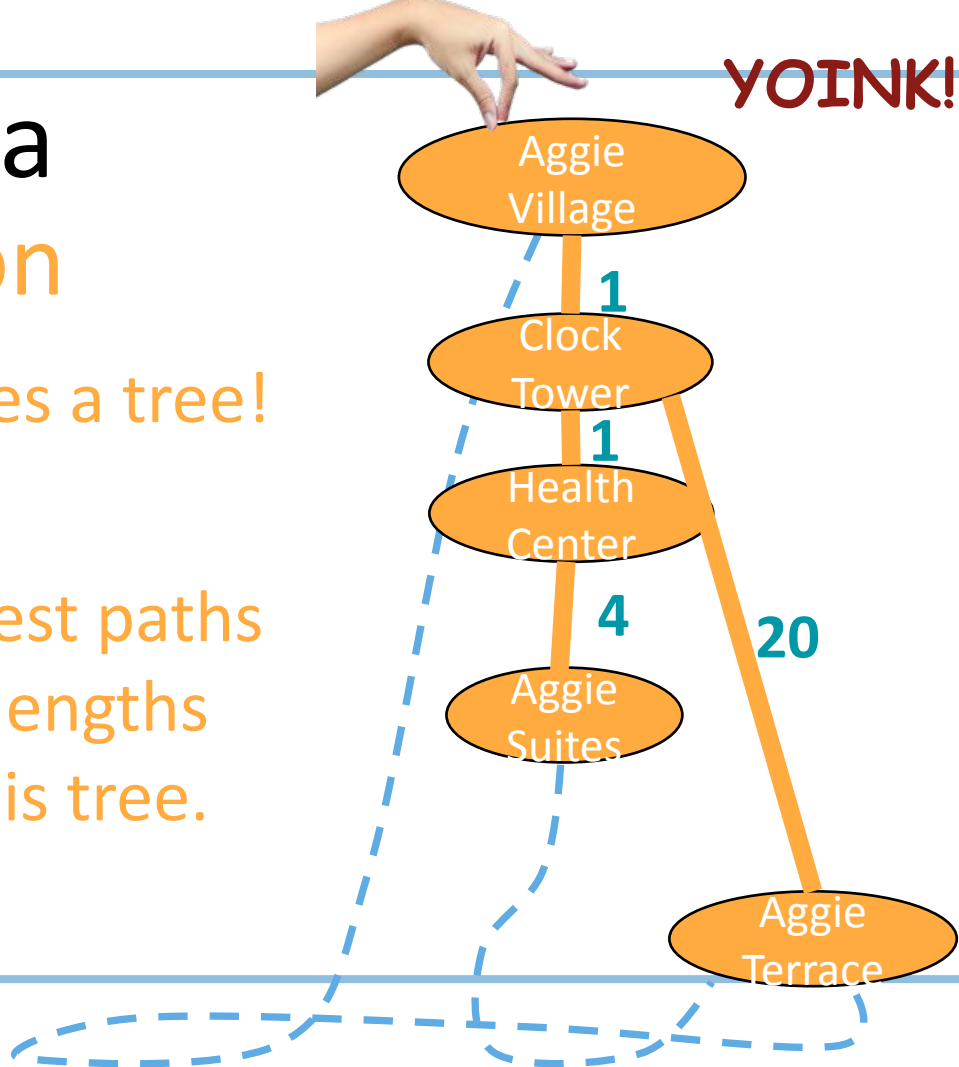
YOINK!



Dijkstra intuition


This creates a tree!

The shortest paths
are the lengths
along this tree.




How do we actually implement this?

- **Without** string and gravity?



Dijkstra by example

How far is a node from Aggie Village?



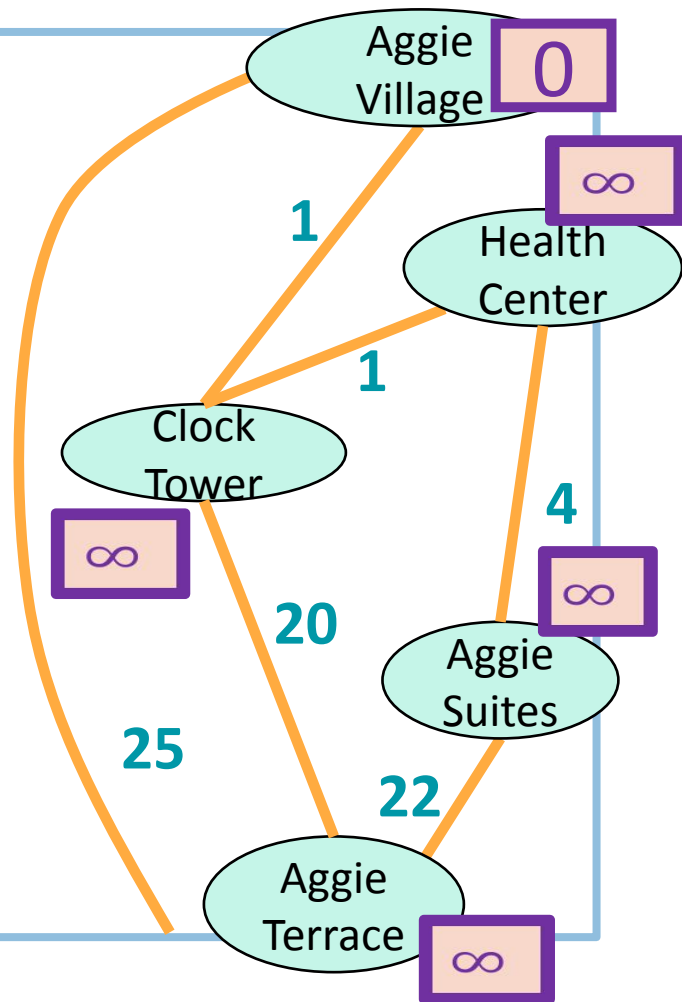
I'm not sure yet



I'm sure



$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet

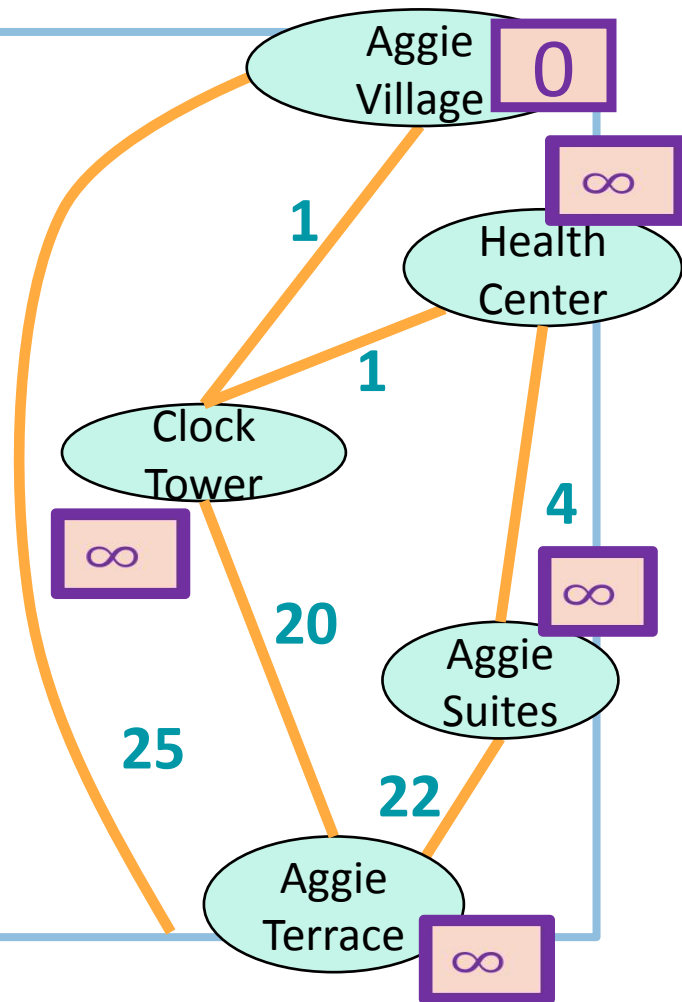


I'm sure




$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure




$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure



$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure



$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**
- Repeat



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure




$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**
- Repeat



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure




$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**
- Repeat



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure




$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**
- Repeat



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure



$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**
- Repeat



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure




$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**
- Repeat



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet

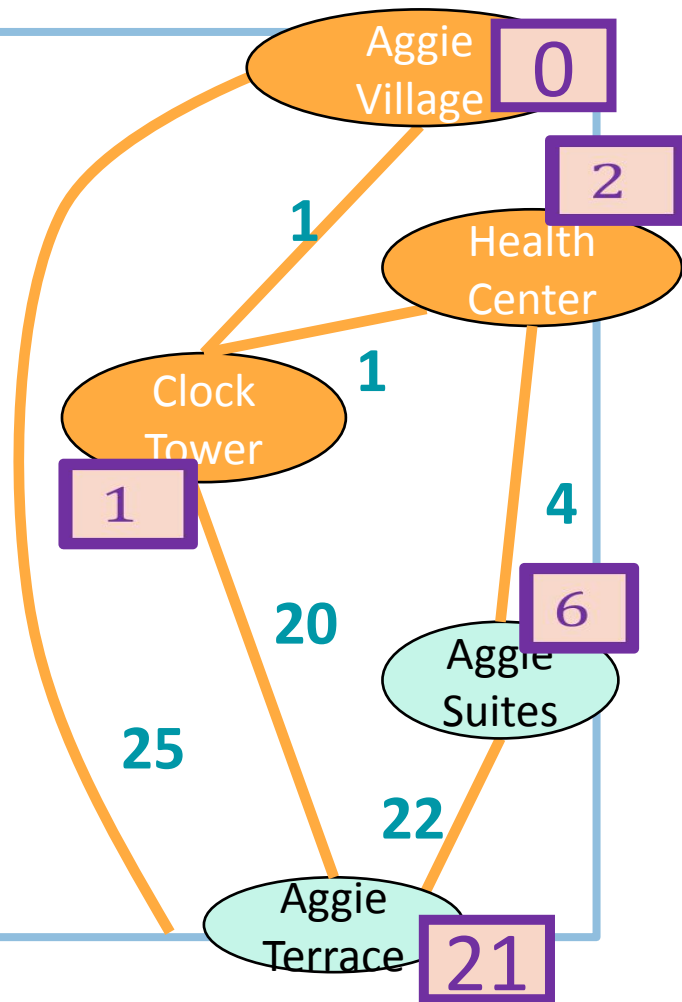


I'm sure




$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**
- Repeat



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure



$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**
- Repeat



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure



$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**
- Repeat



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure




$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
- Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
- Mark u as **sure**
- Repeat



Dijkstra by example

How far is a node from Aggie Village?



I'm not sure yet




I'm sure



$x = d[v]$ is my best **over-estimate** for $\text{dist}(\text{Aggie Village}, v)$.

- Pick the **not-sure** node u with the smallest estimate $d[u]$.
 - Update all u 's neighbors v :
 - $d[v] = \min(d[v], d[u] + \text{edgeWeight}(u,v))$
 - Mark u as **sure**
 - Repeat
- After all nodes are **sure**, $d(\text{Aggie Village}, v) = d[v]$ for all v



Dijkstra's Algorithm

Dijkstra(G, s):

- Set all vertices to **not-sure**
- $d[v] = \infty$ for all v in V
- $d[s] = 0$
- **While** there are **not-sure** nodes:
 - Pick the **not-sure** node u with the smallest estimate $d[u]$.
 - **For** v in u .neighbors:
 - $d[v] \leftarrow \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
 - Mark u as **sure**.
- Now $d(s, v) = d[v]$

Dijkstra's Running time?

Dijkstra(G,s):

- Set all vertices to **not-sure**
- $d[v] = \infty$ for all v in V
- $d[s] = 0$
- **While** there are **not-sure** nodes:
 - Pick the **not-sure** node u with the smallest estimate $d[u]$.
 - **For** v in u .neighbors:
 - $d[v] \leftarrow \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
 - Mark u as **sure**.
- Now $d(s, v) = d[v]$

Dijkstra's Running time?

Dijkstra(G,s):

- Set all vertices to **not-sure**
- $d[v] = \infty$ for all v in V
- $d[s] = 0$
- **While** there are **not-sure** nodes:
 - Pick the **not-sure** node u with the smallest estimate $d[u]$.
 - **For** v in u .neighbors:
 - $d[v] \leftarrow \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
 - Mark u as **sure**.
- Now $d(s, v) = d[v]$
 - n iterations (one per vertex)
 - How long does one iteration take?

Dijkstra's Running time?

Dijkstra(G,s):

- Set all vertices to **not-sure**
- $d[v] = \infty$ for all v in V
- $d[s] = 0$
- **While** there are **not-sure** nodes:
 - Pick the **not-sure** node u with the smallest estimate $d[u]$.
 - **For** v in u .neighbors:
 - $d[v] \leftarrow \min(d[v] , d[u] + \text{edgeWeight}(u,v))$
 - Mark u as **sure**.
- Now $d(s, v) = d[v]$
 - n iterations (one per vertex)
 - How long does one iteration take?

Depends on how we implement it...

We need a data structure that...

Just the inner loop:

- Stores unsure vertices v
- ... And keeps track of $d[v]$
- Can find u with minimum $d[u]$
 - `findMin()`
- Can remove that u
 - `removeMin(u)`
- Can update (decrease) $d[v]$
 - `updateKey(v, d)`

- Pick the **not-sure** node u with the smallest estimate **$d[u]$** .
- Update all u 's neighbors v :
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.

We need a data structure that...

- Stores unsure vertices v
- ... And keeps track of $d[v]$
- Can find u with minimum $d[u]$
 - `findMin()`
- Can remove that u
 - `removeMin(u)`
- Can update (decrease) $d[v]$
 - `updateKey(v, d)`

Just the inner loop:

- Pick the **not-sure** node u with the smallest estimate **$d[u]$** .
- Update all u 's neighbors v :
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.

We need a data structure that...

- Stores unsure vertices v
- ... And keeps track of $d[v]$
- Can find u with minimum $d[u]$
 - `findMin()`
- Can remove that u
 - `removeMin(u)`
- Can update (decrease) $d[v]$
 - `updateKey(v, d)`

Just the inner loop:

- Pick the **not-sure** node u with the smallest estimate **$d[u]$** .
- Update all u 's neighbors v :
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.

What's the runtime...?

How many times are we performing each operation?

We need a data structure that...

- Stores unsure vertices v
- ... And keeps track of $d[v]$
- Can find u with minimum $d[u]$
 - `findMin()`
- Can remove that u
 - `removeMin(u)`
- Can update (decrease) $d[v]$
 - `updateKey(v, d)`

Just the inner loop:

- Pick the **not-sure** node u with the smallest estimate **$d[u]$** .
- Update all u 's neighbors v :
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.

$$= V * (\text{findMin} + \text{removeMin}) + E * \text{updateKey}$$

If we use an array...

- findMin = ???
- removeMin = ???
- updateKey = ???

If we use an array...

- **findMin** = $O(V)$
- **removeMin** = $O(V)$
- **updateKey** = $O(1)$

Running time of Dijkstra

$$= O(V * (\text{findMin} + \text{removeMin}) + E * \text{updateKey})$$

$$= O(V * (V + V) + E * 1)$$

$$= O(V^2 + E)$$

If we use a (balanced) BST...

- findMin = ???
- removeMin = ???
- updateKey = ???

If we use a (balanced) BST...

- **findMin** = $O(\log(V))$
- **removeMin** = $O(\log(V))$
- **updateKey** = $O(\log(V))$

Running time of Dijkstra

$$= O(V * (\text{findMin} + \text{removeMin}) + E * \text{updateKey})$$

$$= O(V * (\log(V) + \log(V)) + E * \log(V))$$

$$= O(V\log(V) + E\log(V))$$

$$= O((V+E)\log(V))$$

Better than an array if the graph is sparse!
aka if m is much smaller than n^2

If we use a hash table...

- findMin = ???
- removeMin = ???
- updateKey = ???

If we use a hash table...

- **findMin** = $O(V)$
- **removeMin** = $O(1)$
- **updateKey** = $O(1)$

Running time of Dijkstra


$$= O(V * (\text{findMin} + \text{removeMin}) + E * \text{updateKey})$$

$$= O(V * (V + 1) + E * 1)$$

$$= O(V^2 + E)$$

Heaps support these operations...

- findMin
- removeMin
- updateKey



- A **min-heap** is a tree-based data structure that has the property that every node has a smaller key than its children.
- A **max-heap** is similar, but every node has a larger key than its children.
- Not covered in this class – see CS166, but we will use them.


Many heap implementations

Nice chart on Wikipedia:

Operation	Binary ^[7]	Leftist	Binomial ^[7]	Fibonacci ^{[7][8]}	Pairing ^[9]	Brodal ^{[10][b]}	Rank-pairing ^[12]	Strict Fibonacci ^[13]
find-min	$\Theta(1)$	$\Theta(1)$	$\Theta(\log n)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$
delete-min	$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(\log n)$	$\mathcal{O}(\log n)^{[c]}$	$\mathcal{O}(\log n)^{[c]}$	$\mathcal{O}(\log n)$	$\mathcal{O}(\log n)^{[c]}$	$\mathcal{O}(\log n)$
insert	$\mathcal{O}(\log n)$	$\Theta(\log n)$	$\Theta(1)^{[c]}$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$
decrease-key	$\Theta(\log n)$	$\Theta(n)$	$\Theta(\log n)$	$\Theta(1)^{[c]}$	$\mathcal{O}(\log n)^{[c][d]}$	$\Theta(1)$	$\Theta(1)^{[c]}$	$\Theta(1)$
merge	$\Theta(n)$	$\Theta(\log n)$	$\mathcal{O}(\log n)^{[e]}$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$


Say we use a binary min-heap...

- findMin = ???
- removeMin = ???
- updateKey = ???



Say we use a binary min-heap...

- **findMin** = $O(1)$
- **removeMin** = $O(\log(V)^*)$
- **updateKey** = $O(\log(V)^*)$



Running time of Dijkstra

$$= O(V * (\text{findMin} + \text{removeMin}) + E * \text{updateKey})$$


$$= O(V * (1 + \log(V)) + E * \log(V))$$

$$= O((V+E)\log(V))$$

*amortized time: any sequence of d `removeMin` calls takes time at most $O(d\log(n))$. But a few of the d may take longer than $O(\log(n))$ and some may take less time..

Say we use a fibonacci min-heap...

- findMin = ???
- removeMin = ???
- updateKey = ???



Say we use a fibonacci heap...

- **findMin** = $O(1)$
- **removeMin** = $O(\log(V)^*)$
- **updateKey** = $O(1)$


Running time of Dijkstra

$$= O(V * (\text{findMin} + \text{removeMin}) + E * \text{updateKey})$$

$$= O(V * (1 + \log(V)) + E * 1)$$


$$= O(V \log(V) + E)$$

*amortized time: any sequence of d **removeMin** calls takes time at most $O(d \log(n))$. But a few of the d may take longer than $O(\log(n))$ and some may take less time..



In practice

Shortest paths on a graph with n vertices and about $5n$ edges




Dijkstra using a Python list to keep track of vertices has quadratic runtime.

Dijkstra using a heap looks a bit more linear (actually $n \log(n)$)

BFS is really fast by comparison! But it doesn't work on weighted graphs.

Dijkstra is used in practice

- eg, OSPF (Open Shortest Path First), a routing protocol for IP networks, uses Dijkstra.



But there are some things it's not so good at...

Dijkstra Drawbacks

- Needs non-negative edge weights.
- If the weights change, we need to re-run the whole thing.
 - in OSPF, a vertex broadcasts any changes to the network, and then every vertex re-runs Dijkstra's algorithm from scratch.

COMP - 285

Analysis of Algorithms

Welcome to COMP 285

Lecture 16: Weighted Graphs, Dijkstra's, A*

Lecturer: Chris Lucas (cflucas@ncat.edu)

