


COMP - 285

Analysis of Algorithms

Welcome to COMP 285

Lecture 17: Minimum Spanning Trees

Lecturer: Chris Lucas (cflucas@ncat.edu)



HW5 due!

Tonight @ 11:59PM ET

HW6 released by EoD!

Due 11/03 @ 11:59PM ET

Final Exam!

Tuesday 12/06 from 2:00pm-4:00pm

Career Office Hours

- Denzel from Meta (sign ups [here](#))
- Resume, general advice, behavioral interview practice, Meta, etc.

Netflix Opportunity!

- Fill out this form ([link](#))
- Check “Meta Classroom”

Quiz!

www.comp285-fall22.ml or Blackboard



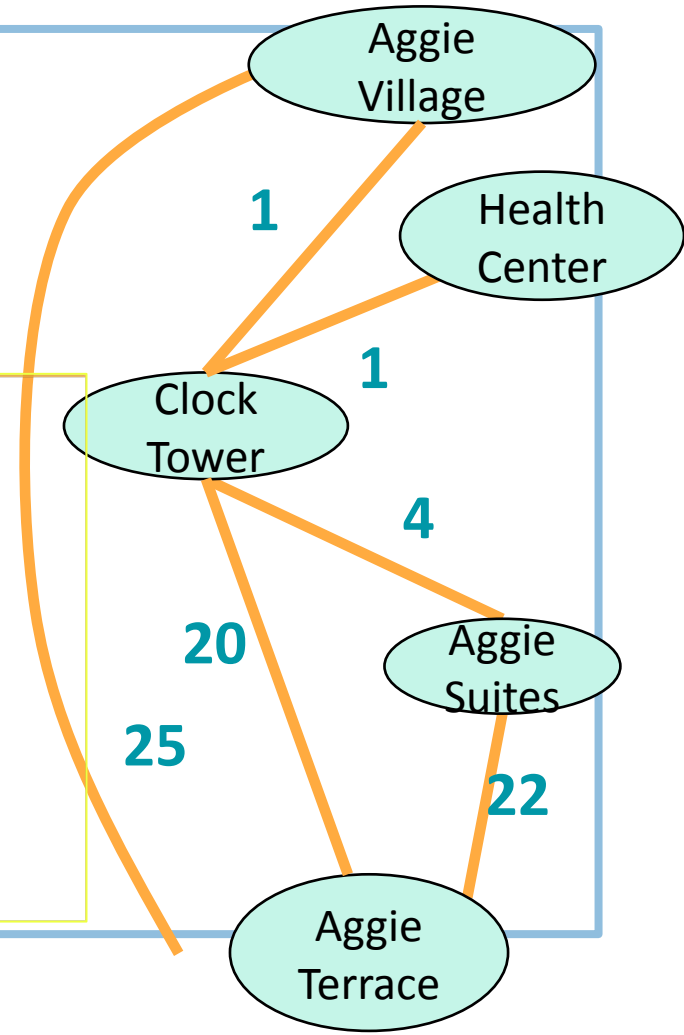
**Recall where we
ended last lecture...**

Dijkstra's algorithm

- Finds shortest paths from Aggie Village to everywhere else.

Dijkstra(G,s):

- Set all vertices to **not-sure**
- $d[v] = \infty$ for all v in V
- $d[s] = 0$
- While** there are **not-sure** nodes:
 - Pick the **not-sure** node u with the smallest estimate $d[u]$.
 - For** v in u .neighbors:
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u,v))$
 - Mark u as **sure**.
- Now $d(s, v) = d[v]$



We need a data structure that...

- Stores unsure vertices v
- ... And keeps track of $d[v]$
- Can find u with minimum $d[u]$
 - `findMin()`
- Can remove that u
 - `removeMin(u)`
- Can update (decrease) $d[v]$
 - `updateKey(v, d)`

Just the inner loop:

- Pick the **not-sure** node u with the smallest estimate **$d[u]$** .
- Update all u 's neighbors v :
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.

We need a data structure that...

- Stores unsure vertices v
- ... And keeps track of $d[v]$
- Can find u with minimum $d[u]$
 - `findMin()`
- Can remove that u
 - `removeMin(u)`
- Can update (decrease) $d[v]$
 - `updateKey(v, d)`

Just the inner loop:

- Pick the **not-sure** node u with the smallest estimate **$d[u]$** .
- Update all u 's neighbors v :
 - $d[v] \leftarrow \min(d[v], d[u] + \text{edgeWeight}(u, v))$
- Mark u as **sure**.

$$= V * (\text{findMin} + \text{removeMin}) + E * \text{updateKey}$$

If we use an array...

- **findMin** = $O(V)$
- **removeMin** = $O(V)$
- **updateKey** = $O(1)$

Running time of Dijkstra

$$= O(V * (\text{findMin} + \text{removeMin}) + E * \text{updateKey})$$

$$= O(V * (V + V) + E * 1)$$

$$= O(V^2 + E)$$

If we use a (balanced) BST...

- **findMin** = $O(\log(V))$
- **removeMin** = $O(\log(V))$
- **updateKey** = $O(\log(V))$

Running time of Dijkstra

$$\begin{aligned} &= O(V * (\text{findMin} + \text{removeMin}) + E * \text{updateKey}) \\ &= O(V * (\log(V) + \log(V)) + E * \log(V)) \\ &= O(V\log(V) + E\log(V)) \\ &= O((V+E)\log(V)) \end{aligned}$$

Say we use a fibonacci heap...

- **findMin** = $O(1)$
- **removeMin** = $O(\log(V)^*)$
- **updateKey** = $O(1)$


Running time of Dijkstra

$$= O(V * (\text{findMin} + \text{removeMin}) + E * \text{updateKey})$$

$$= O(V * (1 + \log(V)) + E * 1)$$

$$= O(V \log(V) + E)$$

*amortized time: any sequence of d **removeMin** calls takes time at most $O(d \log(n))$. But a few of the d may take longer than $O(\log(n))$ and some may take less time..




Big Questions!

- What's a spanning tree/minimum spanning tree?
- How do we find a minimum spanning tree?
- How do we find a minimum spanning tree (pt. 2)?




Big Questions!

- What's a spanning tree/minimum spanning tree?
- How do we find a minimum spanning tree?
- How do we find a minimum spanning tree (pt. 2)?



Say we have an undirected weighted graph...




A **spanning tree** is a **tree** that connects all of the vertices (acyclic). 17

Say we have an undirected weighted graph...

The **cost** of a spanning tree is the sum of the weights on the edges.


This is a spanning tree.

It has cost 67




A **spanning tree** is a **tree** that connects all of the vertices (acyclic). 18

Say we have an undirected weighted graph...



A **spanning tree** is a **tree** that connects all of the vertices (acyclic). 19

What's a Minimum Spanning Tree (MST)?



A **spanning tree** is a **tree** that connects all of the vertices (acyclic). 20

Kahoot!

www.kahoot.it, Code: XXX YYYY

Enter your @aggies.ncat email

What's a Minimum Spanning Tree (MST)?

- True or False? There can only be **one** Spanning Tree in a connected graph G .
 - False
- True or False? There can only be **one** MST in a connected graph G .
 - False
- True or False? In some cases, you may have an MST with $>|V| - 1$ edges.
 - False

Why Minimum Spanning Trees (MSTs)?

- Network Design
 - Connecting cities with roads/electricity/telephone/...
- Cluster Analysis
 - eg, genetic distance
- Image Processing
 - eg, image segmentation
- Useful primitive
 - for other graph algs



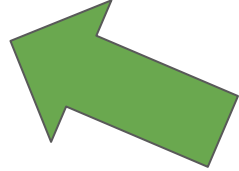


Figure 2: Fully parsimonious minimal spanning tree of 933 SNPs for 282 isolates of *Y. pestis* colored by location. Morelli et al. Nature genetics 2010

Big Questions!


- What's a spanning tree/minimum spanning tree?
- How do we find a minimum spanning tree? 
- How do we find a minimum spanning tree (pt. 2)?



How to find an MST?

- Minimum Spanning Tree (MST) Problem
 - Input: weighted, undirected, connected Graph $G = (V, E)$
 - Output: A Tree $T = (V, E')$, with $E' \subseteq E$ that minimizes the edge weight sum
- This is a formal way of saying “An MST is a Spanning Tree with a connected, weighted graph that has the least total cost when you sum all the edges.”
- Today, we will explore two (greedy) algorithms!

Brainstorm some algorithms...




Back to MSTs


The strategy:

1. Make a **series of choices**, adding edges to the tree.
2. Show that each edge we add is **safe to add**:
 - **We do not create cycles**
3. **Repeat** until we have an MST.


Idea 1: Start growing a tree, (greedily) add the shortest edge we can to grow the tree.




Idea 1: Start growing a tree, (greedily) add the shortest edge we can to grow the tree.




Idea 1: Start growing a tree, (greedily) add the shortest edge we can to grow the tree.




Idea 1: Start growing a tree, (greedily) add the shortest edge we can to grow the tree.




Idea 1: Start growing a tree, (greedily) add the shortest edge we can to grow the tree.




Idea 1: Start growing a tree, (greedily) add the shortest edge we can to grow the tree.




Idea 1: Start growing a tree, (greedily) add the shortest edge we can to grow the tree.



Idea 1: Start growing a tree, (greedily) add the shortest edge we can to grow the tree.



Idea 1: Start growing a tree, (greedily) add the shortest edge we can to grow the tree.




We've Discovered Prim's Algorithm!

- Prim($G = (V, E)$, starting vertex s):
 - Let (s, u) be the lightest edge coming out of s .
 - $MST = \{ (s, u) \}$
 - $verticesVisited = \{ s, u \}$
 - **while** $|verticesVisited| < |V|$:
 - find the lightest edge $\{x, v\}$ in E so that:
 - x is in $verticesVisited$
 - v is not in $verticesVisited$
 - add $\{x, v\}$ to MST
 - add v to $verticesVisited$
 - **return** MST

At most V
iterations of this
while loop.
Time at most E to
go through all the
edges and find the
lightest.


How do we actually implement this?

- Each vertex keeps:
 - the **distance** from itself to the growing spanning tree if you can get there in one edge.
 - how to get there.



How do we actually implement this?

- Each vertex keeps:
 - the **distance** from itself to the growing spanning tree if you can get there in one edge.
 - how to get there.
- Choose the closest vertex, add it.




I'm 7 away.
C is the closest.

I can't get to the
tree in one edge

How do we actually implement this?

- Each vertex keeps:
 - the **distance** from itself to the growing spanning tree if you can get there in one edge.
 - how to get there.**
- Choose the closest vertex, add it.
- Update stored info.



We've Discovered Prim's Algorithm!

algorithm Prim's

Input: Weighted, Undirected, connected Graph $G = (V, E)$ with edge weights w_e

Output: A Tree $T = (V, E')$, with $E' \subseteq E$ that minimizes the edge weight sum

for all $u \in V$:

$\text{cost}(u) = \infty$

$\text{prev}(u) = \text{nil}$

Pick any initial node u_0

$\text{cost}(u_0) = 0$

$H = \text{makequeue}(V)$ (using cost-values as keys)

while H is not empty:

$v = \text{extractMin}(H)$

 for each z of v 's neighbors ($\{v, z\} \in E$):

 if $\text{cost}(z) > w(v, z)$ and $H.\text{contains}(z)$:

$\text{cost}(z) = w(v, z)$

$\text{prev}(z) = v$

$\text{decreasekey}(H, z)$

Prim's Runtime

algorithm Prim's

Input: Weighted, Undirected, connected Graph $G = (V, E)$ with edge weights w_e

Output: A Tree $T = (V, E')$, with $E' \subseteq E$ that minimizes the edge weight sum

for all $u \in V$:

$\text{cost}(u) = \infty$

$\text{prev}(u) = \text{nil}$

Pick any initial node u_0

$\text{cost}(u_0) = 0$

$H = \text{makequeue}(V)$ (using cost-values as keys)

while H is not empty:

$v = \text{extractMin}(H)$

 for each z of v 's neighbors ($\{v, z\} \in E$):

 if $\text{cost}(z) > w(v, z)$ and $H.\text{contains}(z)$:

$\text{cost}(z) = w(v, z)$

$\text{prev}(z) = v$

$\text{decreasekey}(H, z)$

- Time complexity? Let's say extractMin and decreaseKey are on binary heaps ($O(\log(n))$)
- $O(V * \log(V) + E * \log(V))$
 - Because this graph is connected, we know $|V| - 1 \leq |E| < |V|^2$, so we can keep the dominant term and get $O(E * \log(V))$.

Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:




Can't reach x yet
x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:


- Activate the **unreached** vertex u with the **smallest key**.



Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's unreached neighbors v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's unreached neighbors v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to MST.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.



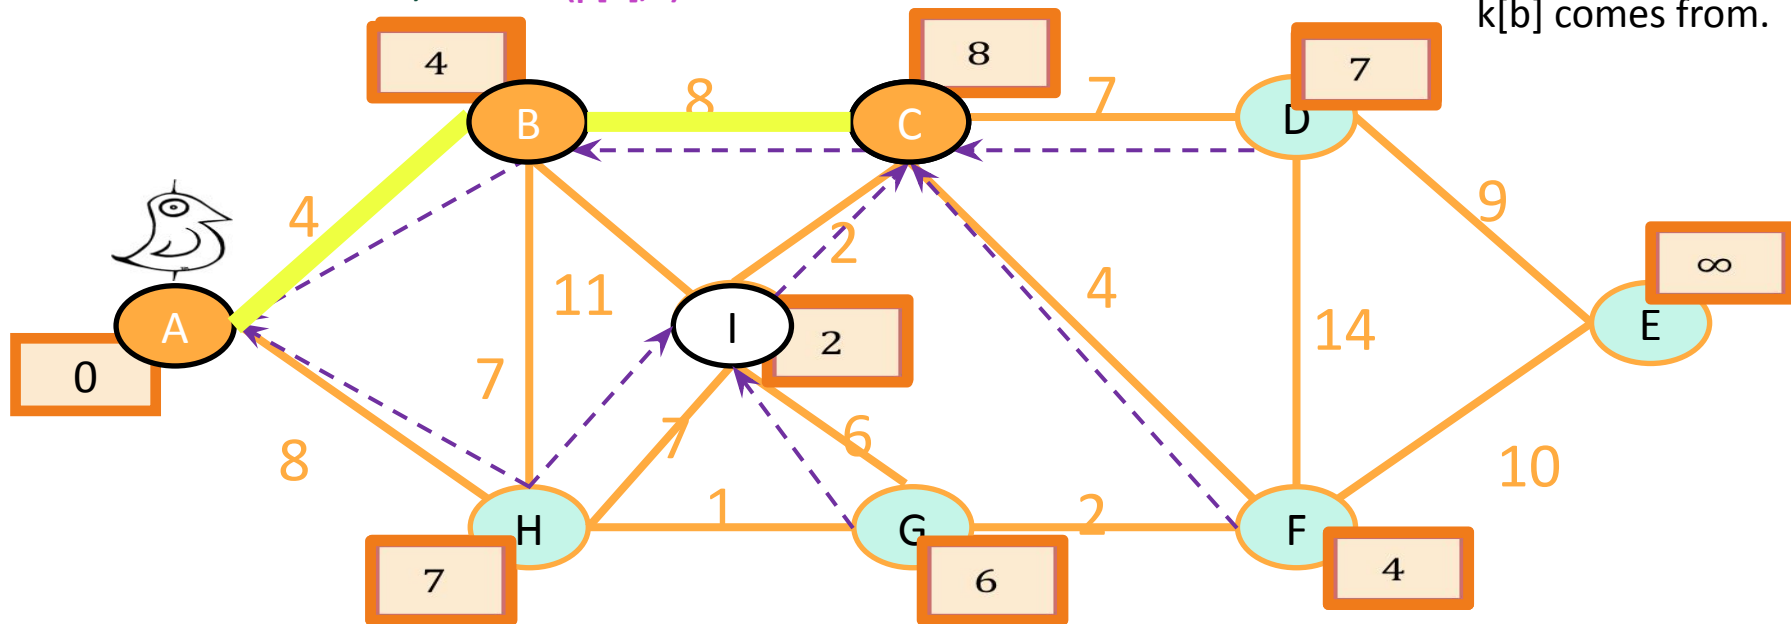
Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.



Efficient implementation

Every vertex has a key and a parent

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u 's **unreached neighbors** v :
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$
- Mark u as **reached**, and **add** $(p[u], u)$ to **MST**.




Can't reach x yet
 x is "active"
Can reach x



$k[x]$ is the distance of x from the growing tree




$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.




This should look pretty familiar...?

- Very similar to Dijkstra's algorithm!
- Differences:
 1. Keep track of $p[v]$ in order to return a tree at the end
 - But Dijkstra's can do that too, that's not a big difference.
 2. Instead of $d[v]$ which we update by
 - $d[v] = \min(d[v], d[u] + w(u,v))$
 - we keep $k[v]$ which we update by
 - $k[v] = \min(k[v], w(u,v))$
- To see the difference, consider:




One thing that is similar: Running time

- Exactly the same as Dijkstra:
 - $O(m \log(n))$ using a balanced tree (as a priority queue).
 - $O(m + n \log(n))$ amortized time if we use a Fibonacci Heap*.




Big Questions!


- What's a spanning tree/minimum spanning tree?
- How do we find a minimum spanning tree?
- How do we find a minimum spanning tree (pt. 2)?




**Idea #2: What if we just always take the cheapest edge?
Whether or not it's connected to what we have so far?**




**Idea #2: What if we just always take the cheapest edge?
Whether or not it's connected to what we have so far?**




**Idea #2: What if we just always take the cheapest edge?
Whether or not it's connected to what we have so far?**




**Idea #2: What if we just always take the cheapest edge?
Whether or not it's connected to what we have so far?**




**Idea #2: What if we just always take the cheapest edge?
Whether or not it's connected to what we have so far?**




**Idea #2: What if we just always take the cheapest edge?
Whether or not it's connected to what we have so far?**



Idea #2: What if we just always take the cheapest edge?
Whether or not it's connected to what we have so far?




**Idea #2: What if we just always take the cheapest edge?
Whether or not it's connected to what we have so far?**




That won't create
a cycle

**Idea #2: What if we just always take the cheapest edge?
Whether or not it's connected to what we have so far?**



That won't create
a cycle

**Idea #2: What if we just always take the cheapest edge?
Whether or not it's connected to what we have so far?**



That won't create
a cycle

We've Discovered Kruskal's Algorithm!

- **KruskalLite($G = (V, E)$):**
 - Sort the edges in E by non-decreasing weight.
 - $MST = \{\}$
 - **for** e in E (in sorted order):
 - **if** adding e to MST won't cause a cycle:
 - add e to MST .
 - **return** MST

E iterations through this loop

How do we check this?

Keep the trees in a special data structure...



“treehouse”?

Union-find data structure (also called disjoint-set data structure)


- Used for storing collections of sets
- Supports:
 - **makeSet(u)**: create a set {u}
 - **find(u)**: return the set that u is in
 - **union(u,v)**: merge the set that u is in with the set that v is in.

`makeSet (x)`

`makeSet (y)`

`makeSet (z)`

`union (x, y)`



Union-find data structure (also called disjoint-set data structure)


- Used for storing collections of sets
- Supports:
 - **makeSet(u)**: create a set {u}
 - **find(u)**: return the set that u is in
 - **union(u,v)**: merge the set that u is in with the set that v is in.

`makeSet (x)`

`makeSet (y)`

`makeSet (z)`


`union (x, y)`



Union-find data structure (also called disjoint-set data structure)

- Used for storing collections of sets
- Supports:
 - **makeSet(u)**: create a set {u}
 - **find(u)**: return the set that u is in
 - **union(u,v)**: merge the set that u is in with the set that v is in.

```
makeSet(x)  
makeSet(y)  
makeSet(z)  
union(x,y)  
find(x)
```



Kruskal Pseudocode (Take 2)


- **Kruskal**($G = (V, E)$):
 - Sort E by weight in non-decreasing order
 - $MST = \{\}$ // initialize an empty tree
 - **for** v in V :
 - **makeSet**(v) // put each vertex in its own tree in the forest
 - **for** (u, v) in E : // go through the edges in sorted order
 - **if** $\text{find}(u) \neq \text{find}(v)$: // if u and v are not in the same tree
 - add (u, v) to MST
 - **union**(u, v) // merge u 's tree with v 's tree
 - **return** MST

At each step of Kruskal's, we maintain a forest (collection of trees).




At each step of Kruskal's, we maintain a forest (collection of trees).

To start, every vertex is in its own tree.




At each step of Kruskal's, we maintain a forest (collection of trees).

Then start merging.



At each step of Kruskal's, we maintain a forest (collection of trees).


Then start merging.



When we add an edge, we merge two trees. We never add an edge within a tree since that would create a cycle.

At each step of Kruskal's, we maintain a forest (collection of trees).


Then start merging.



When we add an edge, we merge two trees. We never add an edge within a tree since that would create a cycle.

At each step of Kruskal's, we maintain a forest (collection of trees).


Then start merging.



When we add an edge, we merge two trees. We never add an edge within a tree since that would create a cycle.

At each step of Kruskal's, we maintain a forest (collection of trees).


Then start merging



When we add an edge, we merge two trees. We never add an edge within a tree since that would create a cycle.

At each step of Kruskal's, we maintain a forest (collection of trees).


Then start merging



When we add an edge, we merge two trees. We never add an edge within a tree since that would create a cycle.

At each step of Kruskal's, we maintain a forest (collection of trees).

Then start merging.




When we add an edge, we merge two trees. We never add an edge within a tree since that would create a cycle.

At each step of Kruskal's, we maintain a forest (collection of trees).

Then start merging.

Stop when we have one big tree!



When we add an edge, we merge two trees. We never add an edge within a tree since that would create a cycle.

○ Worst-case $O(E \log(V))$

Kruskal Pseudocode (Take 2)

- **Kruskal**($G = (V, E)$):
 - Sort E by weight in non-decreasing order
 - $MST = \{\}$ // initialize an empty tree
 - **for** v in V :
 - **makeSet**(v) // put each vertex in its own tree in the forest
 - **for** (u, v) in E : // go through the edges in sorted order
 - **if** $\text{find}(u) \neq \text{find}(v)$: // if u and v are not in the same tree
 - add (u, v) to MST
 - **union**(u, v) // merge u 's tree with v 's tree
 - **return** MST

Compare and contrast

- Prim's

- Looks like Dijkstra's
- Explores nodes based on edges we can reach
- Has a source node
- MST so-far is always connected

Prim might be a better idea
on dense graphs if you can't
radixSort edge weights

- Kruskal's

- Explore edges greedily instead of nodes
- No start node
- MST can be disconnected and slowly become connected as the algorithm runs.

Kruskal might be a better idea
on sparse graphs if you can
radixSort edge weights

COMP - 285

Analysis of Algorithms

Welcome to COMP 285

Lecture 17: Minimum Spanning Trees

Lecturer: Chris Lucas (cflucas@ncat.edu)

